

# R session 2: Customizing R Graphics and Data Manipulations

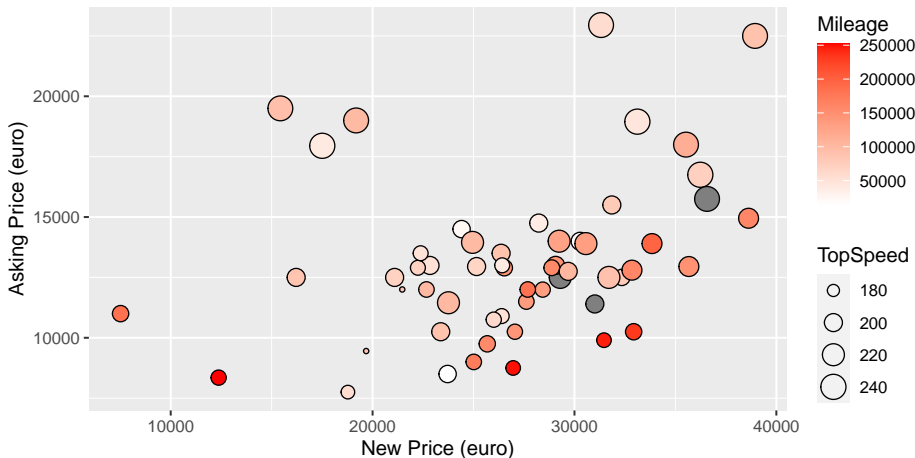
Alina Ferecatu  
Rotterdam School of Management  
Erasmus University

9/24/2020

# Target

## Asking Price versus New Price

56 second-hand VW Golfs from Marktplaats.nl



->After session 2 (+ statistics and data manipulations)

# Today's lecture

- Customizing ggplot2 graphics
- Data manipulations using dplyr

# Recap

# Documents and software

Have the **latest version** of:

- R: <https://CRAN.R-project.org>
- RStudio Desktop: <https://www.rstudio.com>
- Installation instructions on Canvas

# Software requirements

Package tidyverse (for ggplot2 and dplyr), and haven:

```
library("tidyverse")  
library("haven")
```

Optionally, also colourpicker and ggthemes:

```
library("colourpicker")  
library("ggthemes")
```

Data sets: - Prestige.RData: Prestige of occupations in Canada - vwggolf.RData:  
Dutch ads of second-hand 2009 Volkswagen Golfs

# An R session

- Create a new script file: File -> New file -> R Script
- Save script file: Keyboard shortcut: *Ctrl / Cmd + S*
- Do not store the objects (workspace) you created
- Execute the line in which your cursor is with *Ctrl / Cmd + Enter*

# Loading data

File → Open File ... and select the R data file in the dialog

Path relative to the current working directory:

If the file is not in the working directory, specify the full path:

```
load("~/Documents/CMR/session_1/Prestige.RData")
```

→ No need for full path if file is in the working directory

Always use / and not \ (Windows!)



# Basic plotting with package ggplot2

Add together two basic elements:

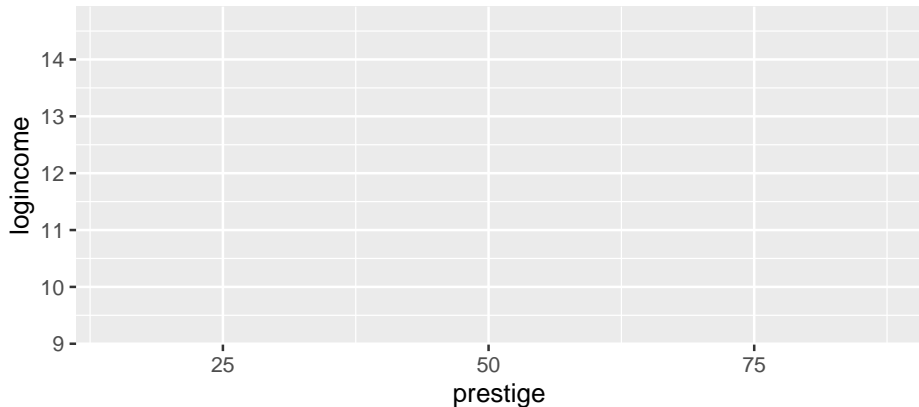
- ❶ Scaffolding defined by `ggplot()`
  - Selects the data set
  - Defines the variables to be used (the aesthetic mapping): function `aes()`
- ❷ Any number of visual representations of the data, known as `geoms`
  - Define the visual representation (the geometric objects): function family `geom_x()`
  - Different elements are added to the plot using the `+` operator

Load the package:

```
library("ggplot2")
```

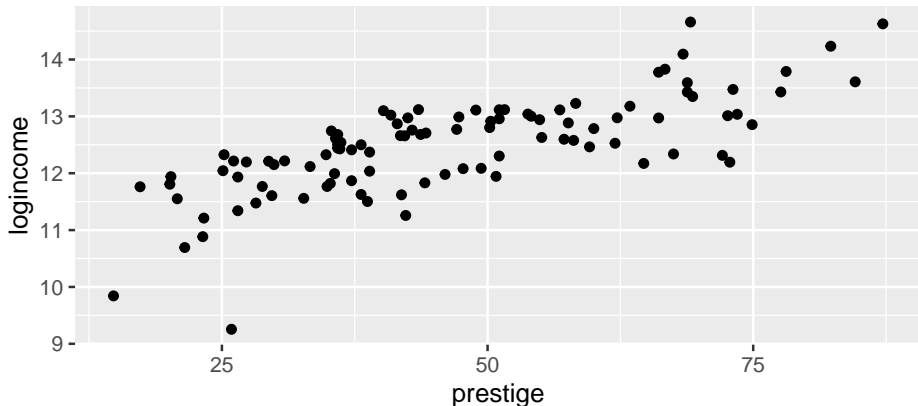
# Scatterplot: Scaffolding

```
ggplot(Prestige, aes(x = prestige, y = logincome))
```



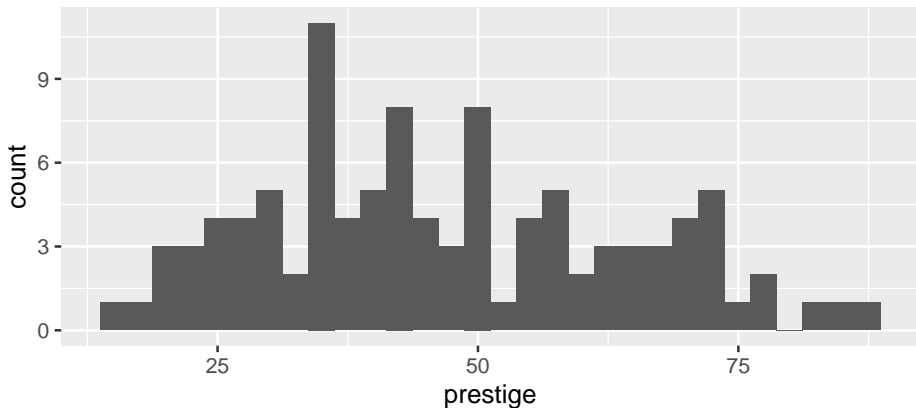
# Scatterplot: Scaffolding + points

```
ggplot(Prestige, aes(x = prestige, y = logincome)) +  
  geom_point()
```



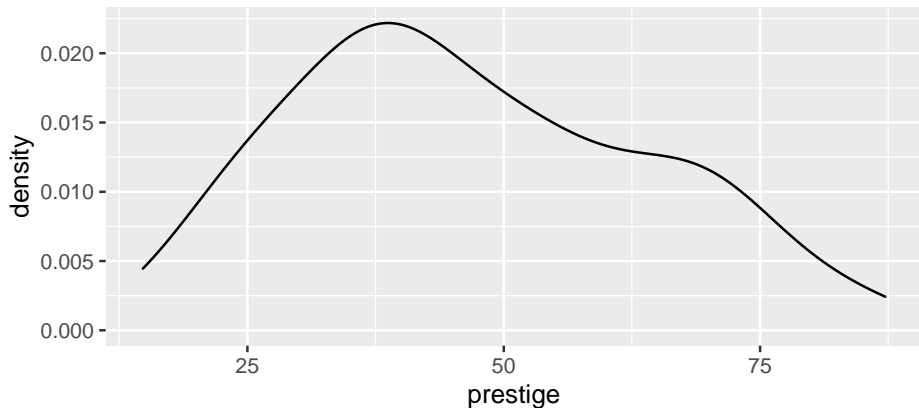
# Histogram

```
ggplot(Prestige, aes(x = prestige)) +  
  geom_histogram()
```



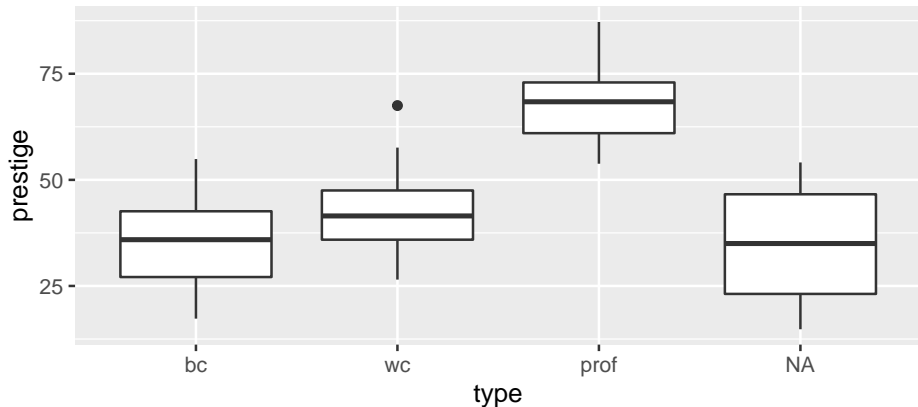
# Density plot

```
ggplot(Prestige, aes(x = prestige)) +  
  geom_density()
```



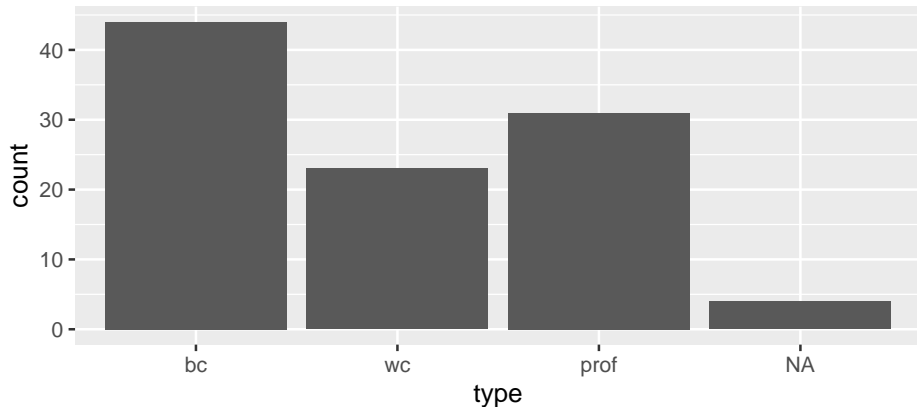
# Conditional boxplot

```
ggplot(Prestige, aes(x = type, y = prestige)) + geom_boxplot()
```



# Barplot

```
ggplot(Prestige, aes(x = type)) + geom_bar()
```



# Some geoms

For a complete list of geoms, click [here](#). Important ones include:

`geom_point()`: Points

`geom_line()`: Lines / time series

`geom_[h/v]line()`: Horizontal or vertical line

`geom_bar()` Bars

`geom_boxplot()` Box and whiskers plot

`geom_density()` Density estimate

`geom_smooth()` Fitted regression line

`geom_text/label` Text

`geom_tile()` Rectangles for heat maps

→ Use appropriate geoms!



# Customizing ggplot2 Graphics

# VW Golf data

Load the data set:

```
load("~/Documents/CMR/session_2/vwgolf.RData")
```

View the data in RStudio:

```
View(vwgolf)
```

Print the dimensions:

```
dim(vwgolf)
```

Univariate summary statistics:

```
summary(vwgolf)
```

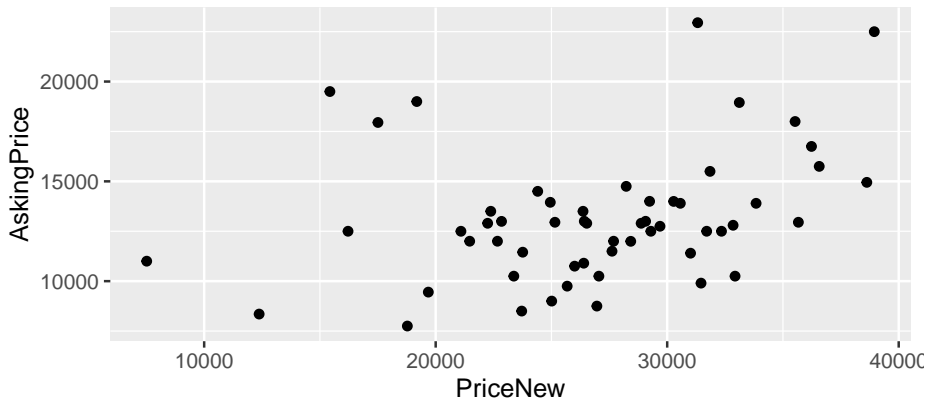
# Exercise 1

Create a new R script, and do the following:

- 1 Create a scatterplot of Mileage against AskingPrice
- 2 Create a scatterplot of Mileage against PriceNew minus AskingPrice
- 3 Create a histogram and density plot of Mileage
- 4 Create boxplots of Mileage conditional on Fuel

# Scatterplot

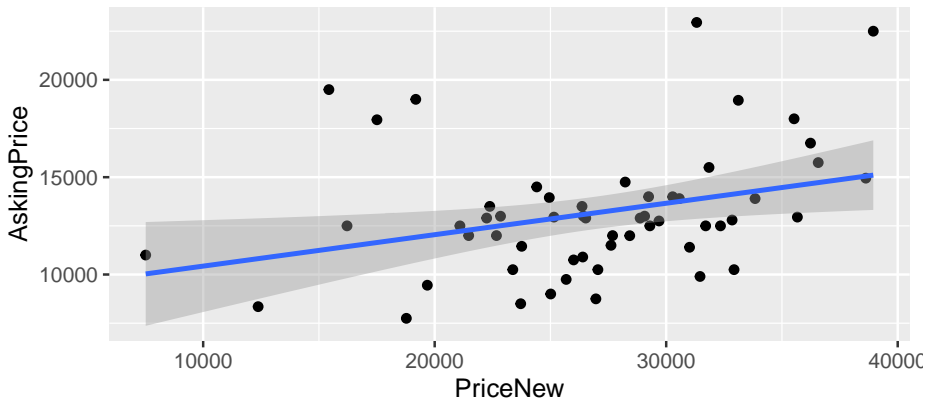
```
ggplot(vwggolf, aes(x = PriceNew, y = AskingPrice)) + geom_point()
```



# Adding multiple geoms

Multiple layers can be added to a plot, such as a regression line:

```
ggplot(vwggolf, aes(x = PriceNew, y = AskingPrice)) +  
  geom_point() + geom_smooth(method = "lm")
```



# Arguments for Fine-tuning

Finetuning can be done for:

**colour** Colour for points/lines (alternatively, color)

**fill** Fill colour for areas

**alpha** Transparency of colours

**shape** Symbol for points

**linetype** Type of line

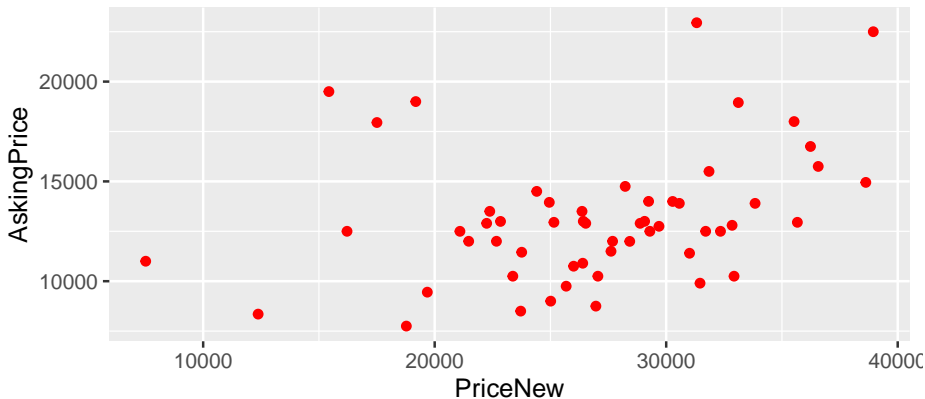
**size** Size of points/lines

If same setting should be used for all points/lines/areas:

—> Best to supply those arguments to `geom_ xxx()`

# Scatterplot with Colour

```
ggplot(vwgolf, aes(x = PriceNew, y = AskingPrice)) +  
geom_point(colour = "red")
```



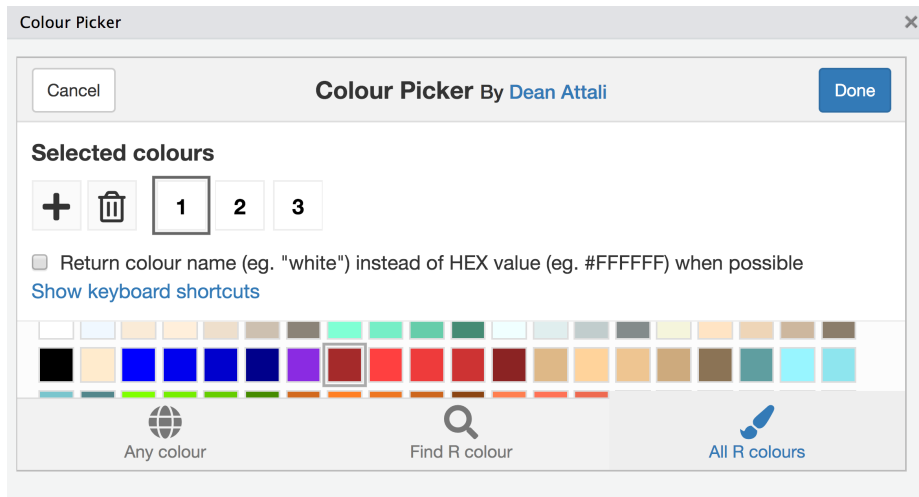
# Colors

```
colours()
```




# Colour Picker

Package colourpicker contains a nice Colour Picker, which will appear in RStudio under Addins after installation:



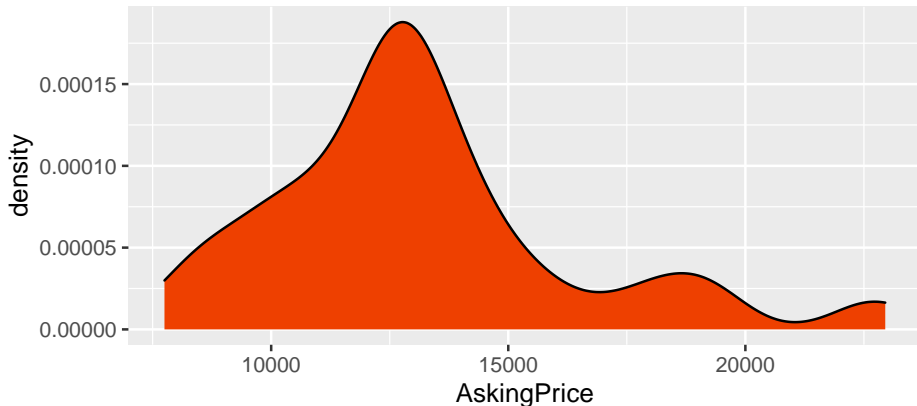
# Scatterplot with Colour

```
ggplot(vwgolf, aes(x = PriceNew, y = AskingPrice)) +  
geom_point(colour = "#00CED1")
```



# Kernel Density Plot with Colour

```
ggplot(vwgf, aes(x = AskingPrice)) +  
  geom_density(fill = "orangered2")
```



# Plot Symbols

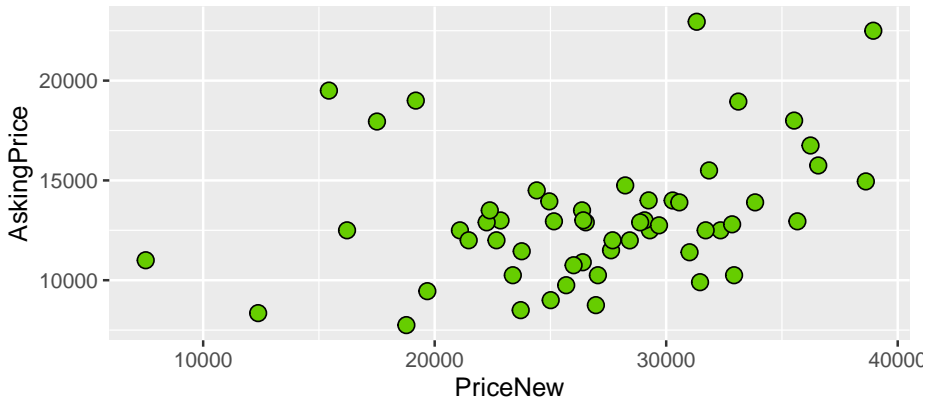


Figure 2: Symbols

—>Specify fill colour for symbols 21—25 with parameter fill

# Adjusted Scatterplot

```
ggplot(vwgolf, aes(x = PriceNew, y = AskingPrice)) +  
  geom_point(colour = "black", fill = "chartreuse3", shape = 21,  
            size = 3)
```



# Linetypes


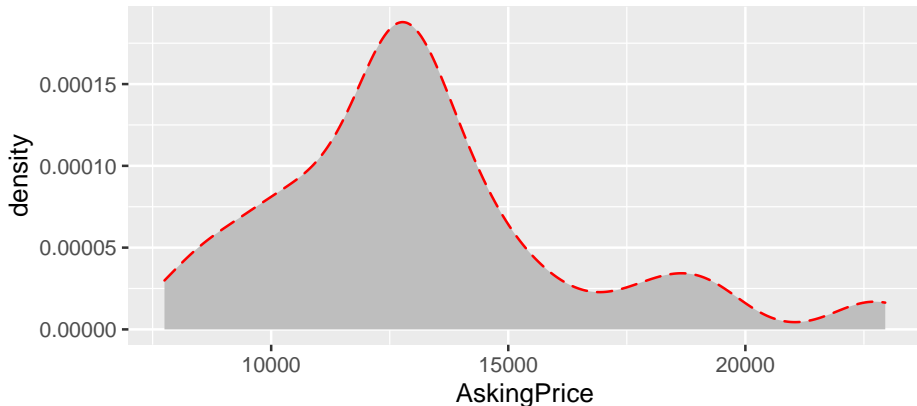


Figure 3: Symbols

—> Specify line type by index or character string

# Kernel Density Plot: Linetype

```
ggplot(vwggolf, aes(x = AskingPrice)) +  
geom_density(fill = "grey", linetype = 5, colour = "red")
```



## Exercise 2

Adjust your plots from the previous exercises as indicated:

- ❶ Create a scatterplot of Mileage against AskingPrice
  - Increase the point size to 3. Set the plotting symbol to 22. Use a purple colour for filling the symbols.
- ❷ Create a histogram of Mileage
  - Change the type of line to dashed, and the colour of the line to red.
  - Also change the fill colour to red. Set the fill colour to be partly transparent using  $\alpha = 0.5$ .



# Mapping Parameters to Aesthetics

Mapping graphical parameters to variables adds more information to a plot Colour / shape / etc. can be mapped to depend on the value of a variable


Specify such mappings in the aesthetic mapping in `aes()`

The scale `xxx()` functions determine how different colours / shapes / etc. are assigned

—> Powerful, but don't add too much information

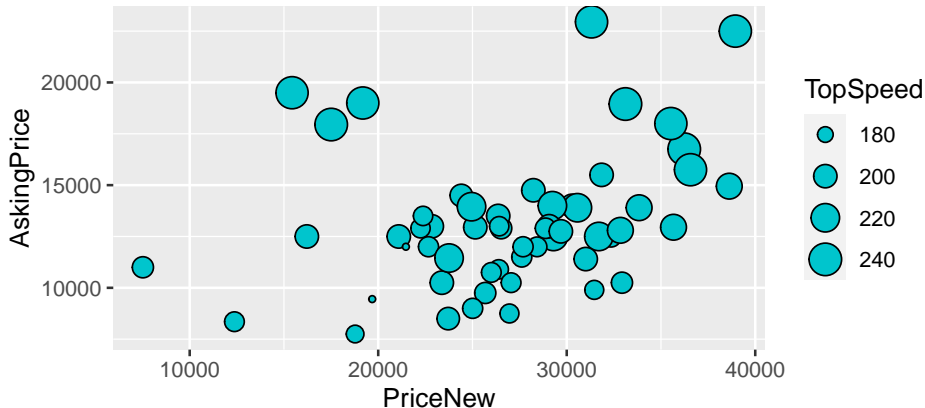
# Scatterplot: Colour Mapping

```
ggplot(vwgolf, aes(x = PriceNew, y = AskingPrice, fill = Fuel)) +  
  geom_point(shape = 21, size = 3)
```



# Scatterplot: Size Mapping

```
ggplot(vwgf,
  aes(x = PriceNew, y = AskingPrice, size = TopSpeed)) +
  geom_point(shape = 21, fill = "turquoise3")
```



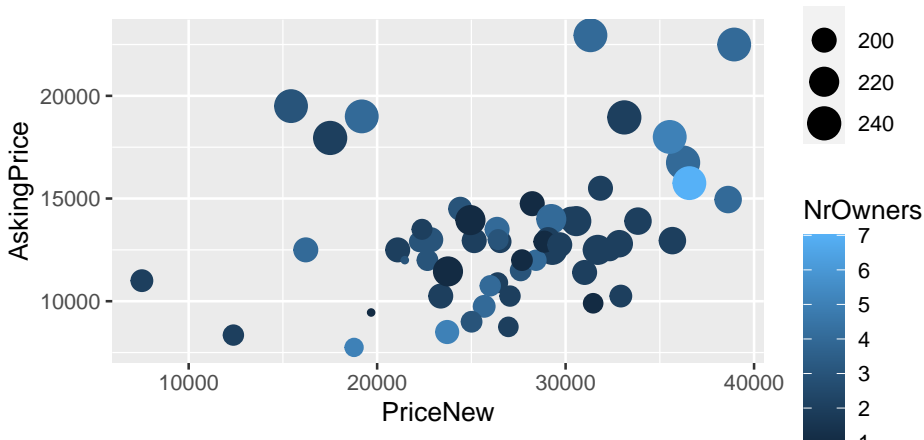
# Scatterplot: Colour Mapping

```
ggplot(vwggolf,  
  aes(x = PriceNew, y = AskingPrice, colour = NrOwners)) +  
  geom_point(size = 4)
```



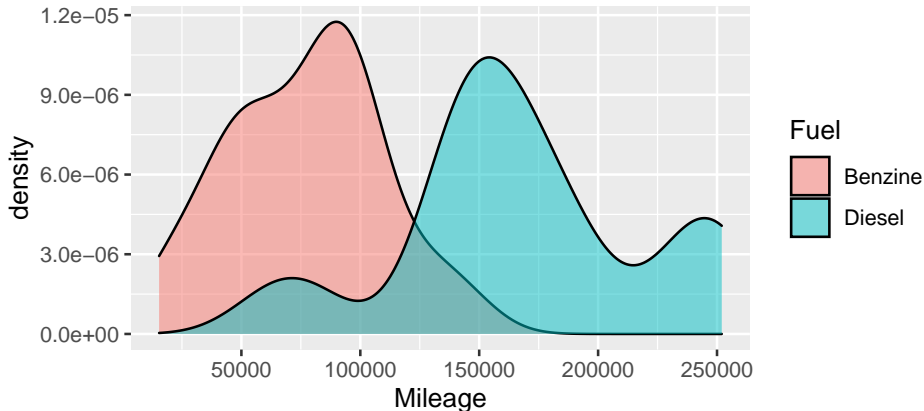
# Scatterplot: Colour and Size Mappings

```
ggplot(vwgf, aes(x = PriceNew, y = AskingPrice,
  colour = NrOwners, size = TopSpeed)) +
  geom_point()
```



# Kernel Density Plot: Fill

```
ggplot(vwgf, aes(x = Mileage, fill = Fuel)) +  
  geom_density(alpha = 0.5)
```



## Exercise 3

- 1 Plot PriceNew against AskingPrice using a scatterplot. Let the size of the points depend on TopSpeed. Use shape 21, and let the fill colour depend on Mileage.
- 2 Create a conditional boxplot of PriceNew conditional on Colour. Map the fill colour of the boxes to Colour.

# Custom Scales

—> Built-in scales are used to obtain default values

Custom discrete scales supplied through function family `scale_xxx_manual()`:

`scale_colour_manual()` - Colour for points/lines

`scale_fill_manual()` - Fill colour for areas

`scale_alpha_manual()` - Transparency of colours

`scale_shape_manual()` - Symbol for points

`scale_linetype_manual()` - Type of line

`scale_size_manual()` - Size of points/lines

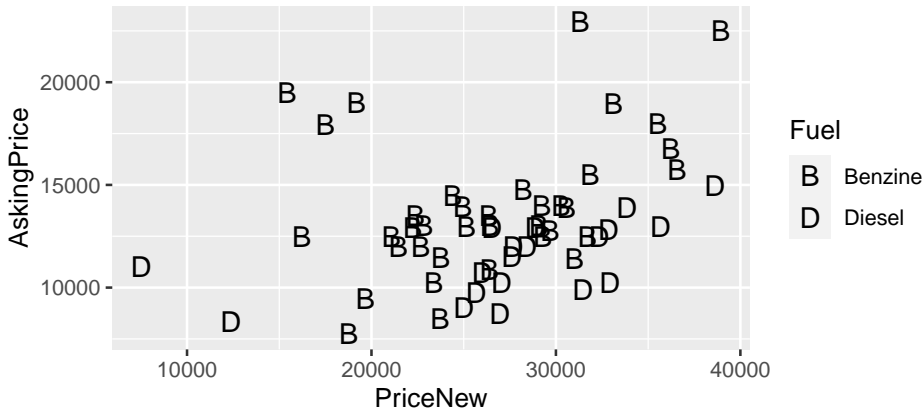
—> Custom values are supplied via argument values

—> Scales are added to a plot with the `+` operator



# Scatterplot: Manual Shapes

```
ggplot(vwgf, aes(x = PriceNew, y = AskingPrice, shape = Fuel)) +  
geom_point(size = 4) + scale_shape_manual(values = c("B", "D"))
```



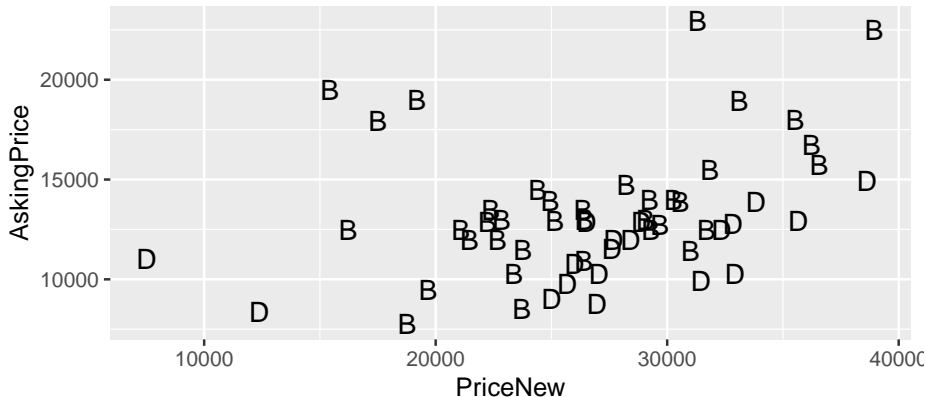
# Removing Legends

For some plots, the legend might not provide new information:

—> Can be removed by setting the argument `show.legend` to `FALSE` in the `geom`

# Scatterplot: Manual Shapes

```
ggplot(vwgfolf, aes(x = PriceNew, y = AskingPrice, shape = Fuel)) +  
  geom_point(size = 4, show.legend = FALSE) +  
  scale_shape_manual(values = c("B", "D"))
```



# Scatterplot: Manual Colours

```
ggplot(vwggolf,  
  aes(x = PriceNew, y = AskingPrice, colour = Imported)) +  
  geom_point() +  
  scale_colour_manual(values = c("#1333E6", "#EB934C"))
```



# Continuous Colour Scales

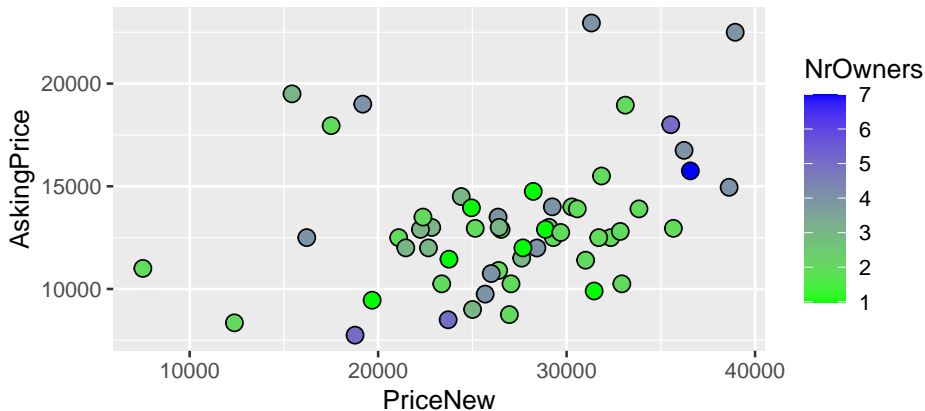
For scales based on continuous variables:

Sequential colour scales are more appropriate

For example, `scale_colour_gradient()` from dark to light

# Scatterplot: Manual Colours

```
ggplot(vwggolf,
  aes(x = PriceNew, y = AskingPrice, fill = NrOwners)) +
  geom_point(shape = 21, size = 3) +
  scale_fill_gradient(low = "green", high = "blue")
```



# Picking Colours

Some colours are better at conveying the information than others

- > Choose carefully, e.g., using packages such as `colorspace`, or the `ggplot2` defaults
- > Never use red and green together

## Exercise 4

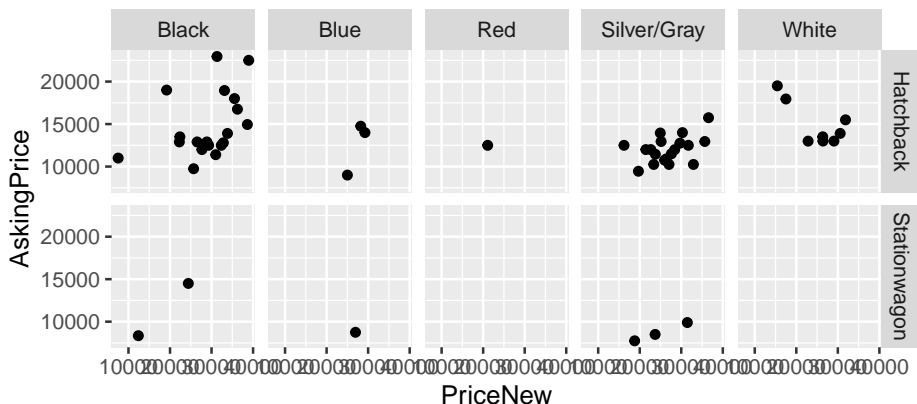
- 1 Create a scatterplot of PriceNew against AskingPrice. Let the point size depend on TopSpeed and the colour depend on Mileage. Use shape 21 with a fill.
- 2 Change the scale of the fill colour to range from white to red. Store the plot as an object named price plot.
- 3 Why are there grey dots in the plot?



# Facets

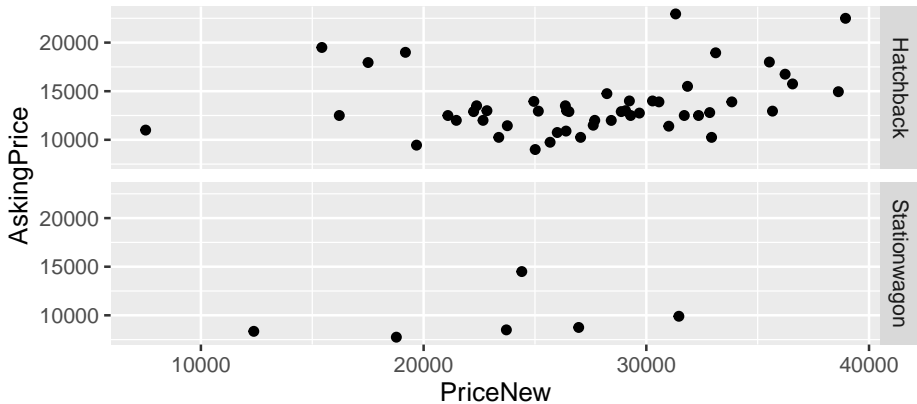
`facet_grid()` constructs multiple panels or facets:

```
ggplot(vwggolf, aes(x = PriceNew, y = AskingPrice)) +  
  geom_point() + facet_grid(Class ~ Colour)
```



# You can use single or multiple variables to define the facets

```
ggplot(vwgolf, aes(x = PriceNew, y = AskingPrice)) +  
geom_point() + facet_grid(Class ~ .)
```



# More Options

Axis limits:

Use `lims()` with arguments `x` and `y`

Titles, subtitles, and axis titles:

Use `labs()` with the arguments `title`, `subtitle`, `x` or `y`

Switch horizontal and vertical axes:

Use `coord_flip()`

Fix the aspect ratio:

Use `coord_fixed()` (Important for representing distances)

—> Simply add to plot object with the `+` operator

## Exercise 5

- 1 Add an appropriate title and subtitle to your price plot from Exercise 4.
- 2 Add neater axis labels.
- 3 Use an aspect ratio of one.

# Themes

The `theme()` function gives control over other aspects of the plot:

—> See `?theme` for an extensive list of options

—> Setting these options makes use of the element `xxx()` functions (such as `element_blank()`)

You can:

Use the default theme, or

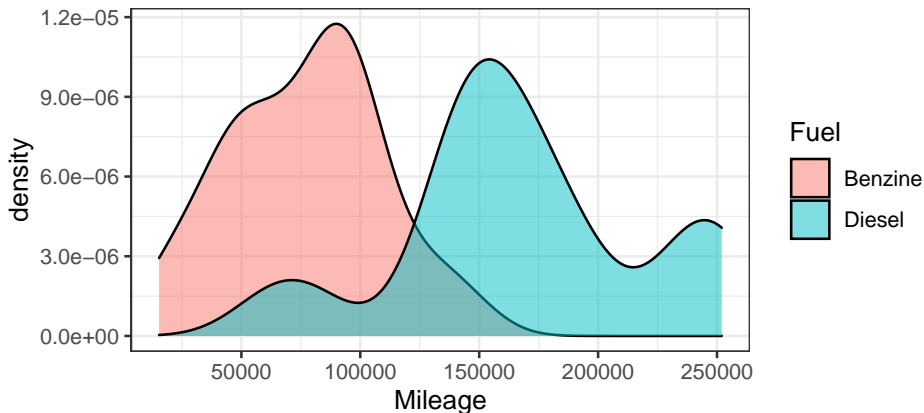
Tweak aspects of the plot to taste, or

Use a built-in theme function, such as `theme_bw()` or `theme_classic()`, or

Even develop your own theme function.

# Printer-friendly theme

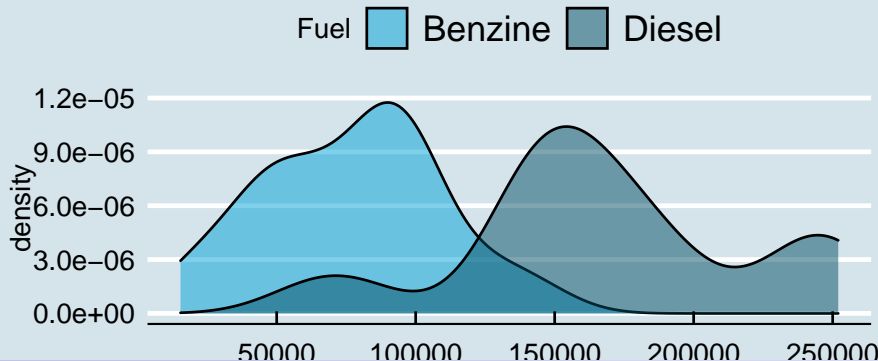
```
ggplot(vwgf, aes(x = Mileage, fill = Fuel)) +  
  geom_density(alpha = 0.5) + theme_bw()
```



# Additional themes

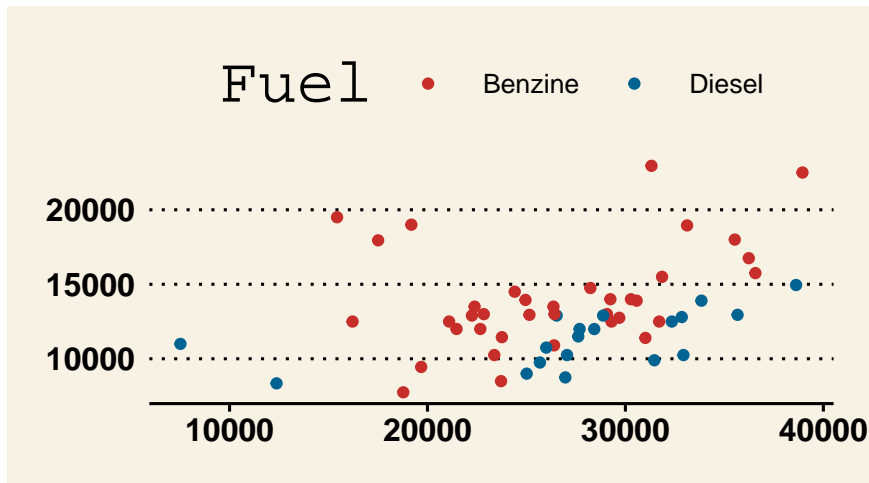
Packages like ggthemes provide additional themes and colour schemes:

```
library(ggthemes)
ggplot(vwggolf, aes(x = Mileage, fill = Fuel)) +
  geom_density(alpha = 0.5) + theme_economist() +
  scale_fill_economist()
```



# Wall Street Journal theme and colours:

```
ggplot(vwggolf, aes(x = PriceNew, y = AskingPrice, colour = Fuel)) +  
geom_point() + theme_ws() + scale_colour_ws()
```





## Aside: Embedding R graphics in documents

- 1 In the Plots tab of the lower right panel, click Export —> Copy to Clipboard. . .
- 2 In the dialog, change width and height if necessary and click *Copy Plot*
- 3 In the document, paste the plot from clipboard

# Data Manipulation with dplyr

# dplyr functions

What to do if the data is not exactly in the form you want?

The dplyr package makes it easy to:

- Pick observations by value with `filter()`
- Reorder the rows using `arrange()`
- Pick variables by name using `select()`
- Create new variables from existing ones using `mutate()`
- Collapse many values to a summary statistic using `summarize()`

# Philosophy

`filter()`, `arrange()`, `select()`, `mutate()` and `summarize()` all work similarly:

- 1 First argument is a data frame
- 2 Subsequent arguments tell R what to do with the data
- 3 The result is a modified data frame

—> All of these can do group-specific computations using `group_by()`

# Subsets of observations

—> Function `filter()` finds observations where a specified condition is true and drops all other observations

Keep only the stationwagons

```
filter(vwgf, Class == "Stationwagon")
```

Keep only the black cars

```
filter(vwgf, Colour == "Black")
```

Keep only the black stationwagons

```
filter(vwgf, Colour == "Black", Class == "Stationwagon")
```

# Comparisons

Operator/ function   **Operation**   Example

`==`   **exactly equal**   `x == y`

`!=`   **not equal**   `x != y`

`<=`   **less or equal**   `x <= y`

`>=`   **greater or equal**   `x >= y`

`<`   **less**   `x < y`

`>`   **greater**   `x > y`

`is.na()`   **is missing**   `is.na(x)`

# Basic logic

Operator/ function   **Operation**   Example

! **not**   !is.na(x)

& **and**   (x > 0) & (x < 1)

| **or**   (x < 0) | (x > 1)

## More filtering rows

Keep only cars with Mileage at most 50 000 km

```
filter(vwgfolf, Mileage <= 50000)
```

—> Note that missing values (NA) are excluded

Display cars with missing values on Mileage:

```
filter(vwgfolf, is.na(Mileage))
```

Display cars with Mileage at most 50 000 km, or with Mileage missing:

```
filter(vwgfolf, Mileage <= 50000 | is.na(Mileage))
```



# Selecting subsets of variables

—> Function `select()` keeps the specified variables (in that order)

Example: Drop all variables after Fuel

These are all equivalent:

```
vw_num <- select(vwggolf, Version, Class, Transmission, Fuel)
vw_num <- select(vwggolf, Version:Fuel)
vw_num <- select(vwggolf, -(Colour:DaysAPK))
head(vw_num)
```

## Rearranging columns, selection based on names

Put AskingPrice and PriceNew first, then everything() else:

```
select(vwgolf, AskingPrice, PriceNew, everything())
```

Select variables with names starting with "Weight"

```
select(vwgolf, starts_with("Weight"))
```

—> Also ends with(), contains(), matches() etc. (see ?starts\_with)

# Sorting rows with arrange()

—> Rows can be sorted based on one or more variables using arrange()

Sort by AskingPrice:

```
arrange(vwggolf, AskingPrice)
```

—> By default, sorted in ascending order

Use desc() to get descending order:

```
arrange(vwggolf, desc(AskingPrice))
```

Using multiple variables:

```
arrange(vwggolf, Fuel, desc(AskingPrice))
```

# Adding or modifying columns

—> Add new or modify existing columns using `mutate()` Add difference between AskingPrice and PriceNew:

```
mutate(vwggolf, PriceDifference = PriceNew - AskingPrice)
```

Transform Mileage to be in 1000's of kilometres:

```
mutate(vwggolf, Mileage = Mileage / 1000)
```

You can do multiple mutations at once:

```
mutate(vwggolf, PriceDifference = PriceNew - AskingPrice,  
Mileage = Mileage / 1000)
```

# Summarise

—> Use `summarize()` to collapse many variables into summary statistics

Specify a name for the summary statistic, and an expression giving the value:

```
summarize(vwgolf,  
  NumberOfCars = length(AskingPrice),  
  MeanPrice = mean(AskingPrice),  
  MinPrice = min(AskingPrice),  
  MaxPrice = max(AskingPrice),  
  SDPrice = sd(AskingPrice))
```

# Conditional computations

- > Use function `group_by()` to make all subsequent computations conditional on one or more grouping variables
- > Use `ungroup()` to undo the grouping behaviour

Summarize AskingPrice per level of Colour:

```
summarize(group_by(vwgf, Colour),  
  NumberOfCars = length(AskingPrice),  
  MeanPrice = mean(AskingPrice),  
  MinPrice = min(AskingPrice),  
  MaxPrice = max(AskingPrice),  
  SDPrice = sd(AskingPrice))
```

# Putting it all together

—> The power of dplyr comes from combining multiple operations in succession

- 1 Retain only diesel cars
- 2 Then select only some variables of interest
- 3 Add a new variable called PriceDifference
- 4 Then summarize by PriceDifference and Mileage

## In one expression:

```
summarize(
  group_by(
    mutate(
      select(
        filter(vwgolf, Fuel == "Diesel"),
        AskingPrice, PriceNew, Mileage, Colour
      ),
      PriceDifference = PriceNew - AskingPrice
    ),
    Colour
  ),
  MeanPriceDifference = mean(PriceDifference),
  MeanMileage = mean(Mileage)
)
```

—> Alternative 1: Store result from intermediate steps

—> Alternative 2: Use the pipe operator `%>%` to chain the commands from left to right



# Chaining with pipes

The function call

```
f(x, y)
```

is the same as

```
x %>% f(y)
```

—>Ctrl/Cmd + Shift + M in RStudio inserts a pipe operator

# Piped data manipulation pipeline

Example revisited:

```
vwgolf %>%  
filter(Fuel == "Diesel") %>%  
select(AskingPrice, PriceNew, Mileage, Colour) %>%  
mutate(PriceDifference = PriceNew - AskingPrice) %>%  
group_by(Colour) %>%  
summarize(MeanPriceDifference = mean(PriceDifference),  
MeanMileage = mean(Mileage))
```

Store the final results for reuse if needed:

```
vwgolf_summary <- vwgolf %>%  
  filter(Fuel == "Diesel") %>%  
  select(AskingPrice, PriceNew, Mileage, Colour) %>%  
  mutate(PriceDifference = PriceNew - AskingPrice) %>%  
  group_by(Colour) %>%  
  summarize(MeanPriceDifference = mean(PriceDifference),  
    MeanMileage = mean(Mileage))
```

# Conclusions

# What we've learned

ggplot2:

- Provides top-class graphics capabilities

- Consistent grammar for lots of plots

- Many options for customization

dplyr:

- Many facilities for constructing recipes for data manipulation

- Useful for exploring and summarizing data

- Useful for preprocessing data before plotting

# Acknowledgements

Thank you Pieter Schoonees and Andreas Alfons  
for contributing to these materials!

# R extra-credit assignment and hackathon

# R extra-credit assignment

Create a new R script (that is, a .R file) in RStudio.

Replicate all analyses posted for the CMR sessions 2, 3, 4, and 5, in the *.html* files.

It is important that your R script contains no errors and is easy to read.

Remember to include code that loads the required packages.

Include comments (lines that start with a *#*) in this text file so that it is clear which code replicates which analysis.

Save your work under a file name containing  
Yourname\_studentNumber\_Rassignment.R

Remember to resave regularly.

Submit the file (one file for all sessions) on Canvas, under *R extra-credit assignment*

Deadline: Oct 8 2020 at 1PM, before the workshop



# R hackathon

Did you spot an error in the code running analyses for the CMR sessions 2, 3, 4, and 5?

- Add a comment starting with `#ERROR` before the analysis
- Propose a correction, and be very specific for which analysis you are proposing the correction

Do you have a more elegant solution to replicate an analysis?

- Add a comment starting with `#SUGGESTION` before the analysis
- Suggest new code for the analysis, and be very specific for which analysis you are making the suggestion

To participate, add your name in the CMR sign-up document, on the sheet named *R hackathon*

# R hackathon winner and prize

*The student reporting most errors and suggestions (#ERROR+#SUGGESTION) wins the R hackathon.*

